



Week 10: *Data Frames*


🏛️ EMSE 4571 / 6571: Intro to Programming for Analytics

👤 John Paul Helveston

📅 March 19, 2026

The Challenger disaster

On January 28, 1986 the space shuttle Challenger exploded



The Challenger disaster

NASA Engineers had the data on temperature & o-ring failure

TEMPERATURE CONCERN ON

SRM JOINTS

27 JAN 1986

HISTORY OF O-RING DAMAGE ON SRM FIELD JOINTS

SRM No.	Erosion Depth (in.)	Perimeter Affected (deg.)	Nominal Dia. (in.)	Top View		Clocking Location (deg.)
				Length of Max Erosion (in.)	Total Heat Affected Length (in.)	
61A LH Center Field**	22A	NONE	0.280	NONE	NONE	36° - 56°
61A LH CENTER FIELD**	22A	NONE	0.280	NONE	NONE	338° - 18°
61C LH Forward Field**	15A	0.010	0.280	4.25	5.25	163
61C RH Center Field (prim)***	15B	0.008	0.280	12.50	58.75	354
61C RH Center Field (sec)***	15B	NONE	0.280	NONE	29.50	354
410 RH Forward Field	13B	0.008	0.280	3.00	NONE	275
41C LH Aft Field*	11A	NONE	0.280	NONE	NONE	--
41B LH Forward Field	10A	0.040	0.280	3.00	14.50	351
STS-2 RH Aft Field	2B	0.053	0.280	--	--	90

*Hot gas path detected in putty. Indication of heat on O-ring, but no damage.

**Soot behind primary O-ring.

***Soot behind primary O-ring, heat affected secondary O-ring.

Clocking location of leak check port - 0 deg.

OTHER SRM-15 FIELD JOINTS HAD NO BLOWHOLES IN PUTTY AND NO SOOT NEAR OR BEYOND THE PRIMARY O-RING.

SRM-22 FORWARD FIELD JOINT HAD PUTTY PATH TO PRIMARY O-RING, BUT NO O-RING EROSION AND NO SOOT BLOWBY. OTHER SRM-22 FIELD JOINTS HAD NO BLOWHOLES IN PUTTY.

BLOW BY HISTORY

SRM-15 WORST BLOW-BY

- 2 CASE JOINTS (50°), (110°) BSC
- MUCH WORSE VIBRALLY THAN SRM-22

SRM-22 BLOW-BY

- 2 CASE JOINTS (30-40°)

SRM-15A, 15, 16A, 18, 23A 24A


- NOZZLE BLOW-BY

HISTORY OF O-RING TEMPERATURES (DEGREES F)


MOTOR	MAT	AMB	O-RING	WIND
DM-4	68	36	47	10 MPH
DM-2	76	45	52	10 MPH
QM-3	72.5	40	48	10 MPH
QM-4	76	48	51	10 MPH
SRM-15	52	64	53	10 MPH
SRM-22	77	78	75	10 MPH
SRM-25	55	26	29	10 MPH
			27	25 MPH

MOTOR O-RING

DM-4	47
DM-2	52
QM-3	48
QM-4	51
SRM-15	53
SRM-22	75
SRM-25	29
	27




History of O-Ring Damage in Field Joints (Cont)



What NASA was shown

Tufte, Edward R. (1997)
Visual Explanations: Images and Quantities, Evidence and Narrative, Graphics Press, Cheshire, Connecticut.

What NASA *should* have been shown



Tufte, Edward R. (1997) *Visual Explanations: Images and Quantities, Evidence and Narrative*, Graphics Press, Cheshire, Connecticut.

"The purpose of computing
is *insight*, not numbers"

- Richard Hamming



Before we start

Make sure you have these packages installed and loaded:

```
install.packages("stringr")  
install.packages("dplyr")  
install.packages("ggplot2")  
install.packages("readr")  
install.packages("here")
```

(At the top of the `practice.R` file)

Remember: you only need to install them once!

Week 10: *Data Frames*

1. Basics

2. Slicing

BREAK

3. External data

Week 10: *Data Frames*


1. Basics

2. Slicing

BREAK

3. External data

The data frame...in Excel



The screenshot shows the Microsoft Excel interface with the 'Data' tab selected. The ribbon includes options for 'From HTML', 'From Text', 'New Database Query', 'Refresh All', 'Connections', 'Properties', 'Edit Links', 'Sort', 'Filter', 'Clear', and 'Advanced'. The formula bar shows 'I17'. The data table is as follows:

	A	B	C	D	E	F	G	H
1	firstName	lastName	instrument	yearOfBirth	deceased			
2	John	Lennon	guitar	1940	TRUE			
3	Paul	McCartney	bass	1942	FALSE			
4	Ringo	Starr	drums	1940	FALSE			
5	George	Harrison	guitar	1943	TRUE			
6								
7								
8								
9								
10								
11								
12								
13								
14								
15								

The data frame...in R

```
beatles <- tibble(  
  firstName = c("John", "Paul", "Ringo", "George"),  
  lastName  = c("Lennon", "McCartney", "Starr", "Harrison"),  
  instrument = c("guitar", "bass", "drums", "guitar"),  
  yearOfBirth = c(1940, 1942, 1940, 1943),  
  deceased   = c(TRUE, FALSE, FALSE, TRUE)  
)
```

beatles

```
#> # A tibble: 4 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar        1940 TRUE  
#> 2 Paul      McCartney  bass          1942 FALSE  
#> 3 Ringo     Starr      drums         1940 FALSE  
#> 4 George    Harrison   guitar        1943 TRUE
```

The data frame...in RStudio

```
View(beatles)
```

	firstName	lastName	instrument	yearOfBirth	deceased
1	John	Lennon	guitar	1940	TRUE
2	Paul	McCartney	bass	1942	FALSE
3	Ringo	Starr	drums	1940	FALSE
4	George	Harrison	guitar	1943	TRUE

Columns: *Vectors* of values (must be same data type)

```
beatles
```

```
#> # A tibble: 4 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar         1940 TRUE  
#> 2 Paul      McCartney  bass           1942 FALSE  
#> 3 Ringo     Starr      drums           1940 FALSE  
#> 4 George    Harrison   guitar          1943 TRUE
```

Extract a column using **\$**

```
beatles$firstName
```

```
#> [1] "John" "Paul" "Ringo" "George"
```

Rows: Information about individual observations

Information about *John Lennon* is in the first row:

```
beatles[1,]
```

```
#> # A tibble: 1 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar          1940 TRUE
```

Information about *Paul McCartney* is in the second row:

```
beatles[2,]
```

```
#> # A tibble: 1 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 Paul      McCartney bass          1942 FALSE
```

Make a data frame with `data.frame()`

```
beatles <- data.frame(  
  firstName = c("John", "Paul", "Ringo", "George"),  
  lastName  = c("Lennon", "McCartney", "Starr", "Harrison"),  
  instrument = c("guitar", "bass", "drums", "guitar"),  
  yearOfBirth = c(1940, 1942, 1940, 1943),  
  deceased   = c(TRUE, FALSE, FALSE, TRUE)  
)
```

beatles

```
#>   firstName  lastName instrument yearOfBirth deceased  
#> 1     John    Lennon    guitar      1940      TRUE  
#> 2     Paul McCartney    bass      1942     FALSE  
#> 3     Ringo     Starr    drums      1940     FALSE  
#> 4     George Harrison    guitar      1943      TRUE
```


Make a data frame with `tibble()`

```
library(dplyr)
```

```
beatles <- tibble(  
  firstName = c("John", "Paul", "Ringo", "George"),  
  lastName  = c("Lennon", "McCartney", "Starr", "Harrison"),  
  instrument = c("guitar", "bass", "drums", "guitar"),  
  yearOfBirth = c(1940, 1942, 1940, 1943),  
  deceased   = c(TRUE, FALSE, FALSE, TRUE)  
)
```

```
beatles
```

```
#> # A tibble: 4 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar        1940 TRUE  
#> 2 Paul      McCartney  bass          1942 FALSE  
#> 3 Ringo     Starr      drums         1940 FALSE  
#> 4 George    Harrison   guitar        1943 TRUE
```

Why I use `tibble()` instead of `data.frame()`

1. The `tibble()` shows the **dimensions** and **data type**.
2. A tibble will only print the first few rows of data when you enter the object name
Example: `faithful` vs. `as_tibble(faithful)`
3. Columns of class `character` are *never* converted into factors (don't worry about this for now...just know that tibbles make life easier with strings).

Note: I use the word "**data frame**" to refer to both `tibble()` and `data.frame()` objects

Data frame vectors must have the same length

```
beatles <- tibble(  
  firstName = c("John", "Paul", "Ringo", "George", "Bob"), # Added "Bob"  
  lastName  = c("Lennon", "McCartney", "Starr", "Harrison"),  
  instrument = c("guitar", "bass", "drums", "guitar"),  
  yearOfBirth = c(1940, 1942, 1940, 1943),  
  deceased   = c(TRUE, FALSE, FALSE, TRUE)  
)
```

```
#> Error in `tibble()`:  
#> ! Tibble columns must have compatible sizes.  
#> • Size 5: Existing data.  
#> • Size 4: Column `lastName`.  
#> i Only values of size one are recycled.
```

Use **NA** for missing values

```
beatles <- tibble(  
  firstName = c("John", "Paul", "Ringo", "George", "Bob"),  
  lastName = c("Lennon", "McCartney", "Starr", "Harrison", NA), # Added NAs  
  instrument = c("guitar", "bass", "drums", "guitar", NA),  
  yearOfBirth = c(1940, 1942, 1940, 1943, NA),  
  deceased = c(TRUE, FALSE, FALSE, TRUE, NA)  
)
```

beatles

```
#> # A tibble: 5 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar         1940 TRUE  
#> 2 Paul      McCartney  bass           1942 FALSE  
#> 3 Ringo     Starr      drums          1940 FALSE  
#> 4 George    Harrison   guitar         1943 TRUE  
#> 5 Bob       <NA>      <NA>           NA NA
```

Dimensions: `nrow()`, `ncol()`, & `dim()`

```
nrow(beatles) # Number of rows
```

```
#> [1] 5
```

```
ncol(beatles) # Number of columns
```

```
#> [1] 5
```

```
dim(beatles) # Number of rows and columns
```

```
#> [1] 5 5
```

Use `names()` or `colnames()` to see the available variables

Get the names of columns:

```
names(beatles) # Usually just use this
```

```
#> [1] "firstName" "lastName" "instrument" "yearOfBirth" "deceased"
```

```
colnames(beatles)
```

```
#> [1] "firstName" "lastName" "instrument" "yearOfBirth" "deceased"
```

Get the names of rows (rarely needed):

```
rownames(beatles)
```

```
#> [1] "1" "2" "3" "4" "5"
```

Changing the column names

Change the column names with `names()` or `colnames()`:

```
names(beatles) <- c('one', 'two', 'three', 'four', 'five')
beatles
```

```
#> # A tibble: 5 × 5
#>   one      two      three  four five
#>   <chr>   <chr>   <chr> <dbl> <lgl>
#> 1 John   Lennon   guitar 1940  TRUE
#> 2 Paul   McCartney bass    1942 FALSE
#> 3 Ringo  Starr    drums  1940 FALSE
#> 4 George Harrison guitar 1943  TRUE
#> 5 Bob    <NA>     <NA>    NA  NA
```

Changing the column names

Make all the column names upper-case:

```
colnames(beatles) <- stringr::str_to_upper(colnames(beatles))  
beatles
```

```
#> # A tibble: 5 × 5  
#>   FIRSTNAME LASTNAME INSTRUMENT YEAROFBIRTH DECEASED  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar         1940 TRUE  
#> 2 Paul      McCartney  bass           1942 FALSE  
#> 3 Ringo     Starr      drums           1940 FALSE  
#> 4 George    Harrison  guitar          1943 TRUE  
#> 5 Bob       <NA>       <NA>           NA NA
```


Combine data frames by columns using `bind_cols()`

Note: `bind_cols()` is from the **dplyr** library

```
names <- tibble(  
  firstName = c("John", "Paul", "Ringo", "George"),  
  lastName  = c("Lennon", "McCartney", "Starr", "Harrison"))  
  
instruments <- tibble(  
  instrument = c("guitar", "bass", "drums", "guitar"))
```

```
bind_cols(names, instruments)
```

```
#> # A tibble: 4 × 3  
#>   firstName lastName instrument  
#>   <chr>      <chr>      <chr>  
#> 1 John      Lennon      guitar  
#> 2 Paul      McCartney  bass  
#> 3 Ringo     Starr      drums  
#> 4 George    Harrison   guitar
```

Combine data frames by rows using `bind_rows()`

Note: `bind_rows()` is from the **dplyr** library

```
members1 <- tibble(  
  firstName = c("John", "Paul"),  
  lastName  = c("Lennon", "McCartney"))  
  
members2 <- tibble(  
  firstName = c("Ringo", "George"),  
  lastName  = c("Starr", "Harrison"))
```

```
bind_rows(members1, members2)
```

```
#> # A tibble: 4 × 2  
#>   firstName lastName  
#>   <chr>      <chr>  
#> 1 John      Lennon  
#> 2 Paul      McCartney  
#> 3 Ringo     Starr  
#> 4 George    Harrison
```

Note: `bind_rows()` requires the **same** columns names:

```
colnames(members2) <- c("firstName", "LastName")  
bind_rows(members1, members2)
```

```
#> # A tibble: 4 × 3  
#>   firstName lastName LastName  
#>   <chr>      <chr>    <chr>  
#> 1 John      Lennon    <NA>  
#> 2 Paul      McCartney <NA>  
#> 3 Ringo     <NA>      Starr  
#> 4 George    <NA>      Harrison
```

Note how `<NA>`s were created

Your turn

06:00

Answer these questions using the `animals_farm` and `animals_pet` data frames:

1. Write code to find how many *rows* are in the `animals_farm` data frame?
2. Write code to find how many *columns* are in the `animals_pet` data frame?
3. Create a new data frame, `animals`, by combining `animals_farm` and `animals_pet`.
4. Change the column names of `animals` to title case.
5. Add a new column to `animals` called `type` that tells if an animal is a `"farm"` or `"pet"` animal.

Week 10: *Data Frames*

1. Basics

2. **Slicing**

BREAK

3. External data

Access data frame columns using the \$ symbol

```
beatles$firstName
```

```
#> [1] "John" "Paul" "Ringo" "George"
```

```
beatles$lastName
```

```
#> [1] "Lennon" "McCartney" "Starr" "Harrison"
```

Creating new variables with the \$ symbol

Add the hometown of the bandmembers:

```
beatles$hometown <- 'Liverpool'  
beatles
```

```
#> # A tibble: 4 × 6  
#>   firstName lastName instrument yearOfBirth deceased hometown  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>      <chr>  
#> 1 John      Lennon      guitar         1940 TRUE       Liverpool  
#> 2 Paul      McCartney  bass           1942 FALSE      Liverpool  
#> 3 Ringo     Starr      drums          1940 FALSE      Liverpool  
#> 4 George    Harrison   guitar         1943 TRUE       Liverpool
```

Creating new variables with the \$ symbol

Add a new `alive` variable:

```
beatles$alive <- c(FALSE, TRUE, TRUE, FALSE)
beatles
```

```
#> # A tibble: 4 × 7
#>   firstName lastName instrument yearOfBirth deceased hometown alive
#>   <chr>      <chr>      <chr>          <dbl> <lgl>      <chr>    <lgl>
#> 1 John      Lennon      guitar          1940 TRUE      Liverpool FALSE
#> 2 Paul      McCartney bass            1942 FALSE     Liverpool TRUE
#> 3 Ringo     Starr       drums            1940 FALSE     Liverpool TRUE
#> 4 George    Harrison    guitar            1943 TRUE      Liverpool FALSE
```


You can compute new variables from current ones

Compute and add the age of the bandmembers:

```
beatles$age <- 2023 - beatles$yearOfBirth  
beatles
```

```
#> # A tibble: 4 × 8  
#>   firstName lastName instrument yearOfBirth deceased hometown alive age  
#>   <chr>      <chr>      <chr>      <dbl> <lgl>      <chr>      <lgl> <dbl>  
#> 1 John      Lennon      guitar      1940 TRUE      Liverpool FALSE 83  
#> 2 Paul      McCartney bass        1942 FALSE     Liverpool TRUE 81  
#> 3 Ringo     Starr      drums        1940 FALSE     Liverpool TRUE 83  
#> 4 George    Harrison   guitar        1943 TRUE      Liverpool FALSE 80
```

Access elements by index: `df[row, column]`

General form for indexing elements:

```
df[row, column]
```

Select the element in row 1, column 2:

```
beatles[1, 2]
```

```
#> # A tibble: 1 × 1  
#>   lastName  
#>   <chr>  
#> 1 Lennon
```

Select the elements in rows 1 & 2 and columns 2 & 3:

```
beatles[c(1, 2), c(2, 3)]
```

```
#> # A tibble: 2 × 2  
#>   lastName instrument  
#>   <chr>      <chr>  
#> 1 Lennon    guitar  
#> 2 McCartney bass
```

Leave row or column "blank" to select all

```
beatles[c(1, 2),] # Selects all COLUMNS for rows 1 & 2
```

```
#> # A tibble: 2 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 John      Lennon      guitar         1940 TRUE  
#> 2 Paul      McCartney bass           1942 FALSE
```

```
beatles[,c(1, 2)] # Selects all ROWS for columns 1 & 2
```

```
#> # A tibble: 4 × 2  
#>   firstName lastName  
#>   <chr>      <chr>  
#> 1 John      Lennon  
#> 2 Paul      McCartney  
#> 3 Ringo     Starr  
#> 4 George    Harrison
```

Negative indices exclude row / column

```
beatles[-1, ] # Select all ROWS except the first
```

```
#> # A tibble: 3 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 Paul      McCartney bass           1942 FALSE  
#> 2 Ringo     Starr      drums          1940 FALSE  
#> 3 George    Harrison guitar         1943 TRUE
```

```
beatles[, -1] # Select all COLUMNS except the first
```

```
#> # A tibble: 4 × 4  
#>   lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>          <dbl> <lgl>  
#> 1 Lennon    guitar         1940 TRUE  
#> 2 McCartney bass           1942 FALSE  
#> 3 Starr     drums          1940 FALSE  
#> 4 Harrison  guitar         1943 TRUE
```

You can select columns by their names

Note: you don't need the comma to select an entire column

One column

```
beatles['firstName']
```

```
#> # A tibble: 4 × 1
#>   firstName
#>   <chr>
#> 1 John
#> 2 Paul
#> 3 Ringo
#> 4 George
```

Multiple columns

```
beatles[c('firstName', 'lastName')]
```

```
#> # A tibble: 4 × 2
#>   firstName lastName
#>   <chr>      <chr>
#> 1 John      Lennon
#> 2 Paul      McCartney
#> 3 Ringo     Starr
#> 4 George    Harrison
```

Use logical indices to *filter* rows

Which Beatles members are still alive?

Create a logical vector using the `deceased` column:

```
beatles$deceased == FALSE
```

```
#> [1] FALSE TRUE TRUE FALSE
```

Insert this logical vector in the ROW position of `beatles[, ]`:

```
beatles[beatles$deceased == FALSE,]
```

```
#> # A tibble: 2 × 5  
#>   firstName lastName instrument yearOfBirth deceased  
#>   <chr>      <chr>      <chr>          <dbl> <lgl>  
#> 1 Paul      McCartney bass           1942 FALSE  
#> 2 Ringo     Starr      drums           1940 FALSE
```

Your turn

10:00

Answer these questions using the `beatles` data frame:

1. Create a new column, `playsGuitar`, which is `TRUE` if the band member plays the guitar and `FALSE` otherwise.
2. Filter the data frame to select only the rows for the band members who have four-letter first names.
3. Create a new column, `fullName`, which contains the band member's first and last name separated by a space (e.g. `"John Lennon"`)

Intermission

05 : 00

Week 10: *Data Frames*

1. Basics

2. Slicing

BREAK

3. External data

Getting data into R

Options:

1. Load external packages
2. Read in external files (usually a `.csv*` file)

*csv = "comma-separated values"

Data from an R package

```
library(ggplot2)
```

See which data frames are available in a package:

```
data(package = "ggplot2")
```

Find out about package data sets with ?

```
?msleep
```

```
msleep {ggplot2}
```

```
An updated and expanded version of the mammals sleep dataset
```

```
Description
```

```
This is an updated and expanded version of the mammals sleep dataset. Updated sleep times
```

Previewing data frames: `msleep`

Look at the data in a "spreadsheet"-like way:

```
View(msleep)
```

This is "read-only" so you can't corrupt the data 😊

My favorite quick summary: `glimpse()`

Preview each variable with `str()` or `glimpse()`

```
library(dplyr)

glimpse(msleep)
```

```
#> Rows: 83
#> Columns: 11
#> $ name      <chr> "Cheetah", "Owl monkey", "Mountain beaver", "Greater short-t
#> $ genus     <chr> "Acinonyx", "Aotus", "Aplodontia", "Blarina", "Bos", "Bradyp
#> $ vore      <chr> "carni", "omni", "herbi", "omni", "herbi", "herbi", "carni",
#> $ order     <chr> "Carnivora", "Primates", "Rodentia", "Soricomorpha", "Artiod
#> $ conservation <chr> "lc", NA, "nt", "lc", "domesticated", NA, "vu", NA, "domesti
#> $ sleep_total <dbl> 12.1, 17.0, 14.4, 14.9, 4.0, 14.4, 8.7, 7.0, 10.1, 3.0, 5.3,
#> $ sleep_rem  <dbl> NA, 1.8, 2.4, 2.3, 0.7, 2.2, 1.4, NA, 2.9, NA, 0.6, 0.8, 0.7
#> $ sleep_cycle <dbl> NA, NA, NA, 0.1333333, 0.6666667, 0.7666667, 0.3833333, NA,
#> $ awake     <dbl> 11.90, 7.00, 9.60, 9.10, 20.00, 9.60, 15.30, 17.00, 13.90, 2
#> $ brainwt   <dbl> NA, 0.01550, NA, 0.00029, 0.42300, NA, NA, NA, 0.07000, 0.09
#> $ bodywt    <dbl> 50.000, 0.480, 1.350, 0.019, 600.000, 3.850, 20.490, 0.045,
```

Also very useful for quick checks: `head()` and `tail()`

View the **first** 6 rows with `head()`

```
head(msleep)
```

```
#> # A tibble: 6 × 11
#>   name          genus
#>   <chr>         <chr>
#> 1 Cheetah       Acinonyx
#> 2 Owl monkey    Aotus
#> 3 Mountain beaver Aplodontia
#> 4 Greater short-tailed shrew Blarina
#> 5 Cow           Bos
#> 6 Three-toed sloth Bradypus
```

View the **last** 6 rows with `tail()`

```
tail(msleep)
```

```
#> # A tibble: 6 × 11
#>   name          genus     vore
#>   <chr>         <chr>   <chr>
#> 1 Tenrec        Tenrec  omni
#> 2 Tree shrew    Tupaia  omni
#> 3 Bottle-nosed dolphin Tursiops carni
#> 4 Genet         Genetta carni
#> 5 Arctic fox    Vulpes  carni
#> 6 Red fox       Vulpes  carni
```

Importing an external data file

Note the `data.csv` file in your `data` folder.

- **DO NOT** double-click it!
- **DO NOT** open it in Excel!

Excel can **corrupt** your data!

(Don't believe me? read [this](#))

If you **must** open it in Excel:

- Make a copy
- Open the copy

Steps to importing external data files

1. Create a path to the data

```
library(here)  
pathToData <- here('data', 'data.csv')  
pathToData
```

```
#> [1] "/Users/jhelvy/gh/teaching/P4A/2026-Spring/class/8-data-frames/data/data.csv"
```

2. Import the data

```
library(readr)  
df <- read_csv(pathToData)
```

Using the **here** package to make file paths

The `here()` function builds the path to your **root** to your *working directory* (this is where your `.Rproj` file lives!)

```
here()
```

```
#> [1] "/Users/jhelvy/gh/teaching/P4A/2026-Spring/class/8-data-frames"
```

The `here()` function builds the path to files *inside* your working directory

```
path_to_data <- here('data', 'data.csv')  
path_to_data
```

```
#> [1] "/Users/jhelvy/gh/teaching/P4A/2026-Spring/class/8-data-frames/data/data.csv"
```

Avoid hard-coding file paths!


(they can break on different computers)

```
path_to_data <- 'data/data.csv'  
path_to_data
```

```
#> [1] "data/data.csv"
```



Use the **here** package
to make file paths



Art by [Allison Horst](#)

Use `read_csv()`, not `read.csv()`



`read.csv()`



`read_csv()`

```
path_to_data <- here('data', 'data.csv')  
data <- read_csv(path_to_data)
```

Save data frame as csv with `write_csv()`

1. Create a path to where you want to save the data

```
library(here)
pathToData <- here('data', 'data_new.csv')
pathToData
```

```
#> [1] "/Users/jhelvy/gh/teaching/P4A/2026-Spring/class/8-data-frames/data/data_new.csv"
```

2. Write the file to disc

```
library(readr)
write_csv(data, pathToData)
```

Your turn

10:00

- 1) Use the `here()` and `read_csv()` functions to load the `data.csv` file that is in the `data` folder. Name the data frame object `df`.
- 2) Use the `df` object to answer the following questions:
 - How many rows and columns are in the data frame?
 - Preview the different columns. What do you think this data is about? What might one row represent? What type of data is each column? (don't need to type this out...just inspect the data)
 - How many unique airports are in the data frame?
 - What is the earliest and latest observation in the data frame?
 - What is the lowest and highest cost of any one repair in the data frame?
- 3) Create a subset of the data called `data_dc` that contains only observations from DC-area airports (IAD, DCA, & BWI), then use `write_csv()` to save it in your "data" folder as `data_cs.csv`.

Next week: better data wrangling with **dplyr**



Art by [Allison Horst](#)

Select rows with `filter()`

Example: Filter rows to find which Beatles members are still alive?

Base R:

```
beatles[beatles$deceased == FALSE,]
```

dplyr:


```
filter(beatles, deceased == FALSE)
```

In 2 weeks: plotting with **ggplot2**

Translate *data*...

...into *information*

```
#> # A tibble: 11 × 2  
#>   brainwt  bodywt  
#>   <dbl>   <dbl>  
#> 1 0.001    0.06  
#> 2 0.0066    1  
#> 3 0.00014  0.005  
#> 4 0.0108    3.5  
#> 5 0.0123    2.95  
#> 6 0.0063    1.7  
#> 7 4.60     2547  
#> 8 0.0003    0.023  
#> 9 0.655     521  
#> 10 0.419    187  
#> 11 0.0035    0.77
```



A note about HW 8

- You have what you need to start now.
- It will be *much* easier if you use the **dplyr** functions (i.e. read ahead). Feel free to try using them on this assignment.

Extra Practice!

10:00

1. Install the **dslabs** package.
2. Load the package with `library(dslabs)`
3. Use `data(package = "dslabs")` to see the different data sets in this package.
4. Pick one.
5. Answer these questions about the dataset you chose:
 - What is the dataset about?
 - How many observations are in the data frame?
 - What is the original source of the data?
 - What type of data is each variable?
 - Find one thing interesting about it to share.